




Predicting Lung Disease Recovery



COMP 4448: DS Tools 2
Logan Hahn



Purpose & Significance

- Employ and compare three classification models to determine which could more accurately predict whether a lung disease patient will recover.
- Predicting recovery outcomes in lung disease patients can provide a massive value to both the patients and the healthcare providers.

Research Question

Which algorithm, Logistic Regression, Random Forest, or XGBoost, would be the most effective in predicting whether or not a patient with lung disease will recover based on their demographic, lifestyle, and clinical characteristics?



Dataset

Source: <https://www.kaggle.com/datasets/samikshadalvi/lungs-diseases-dataset>

Input Variables:

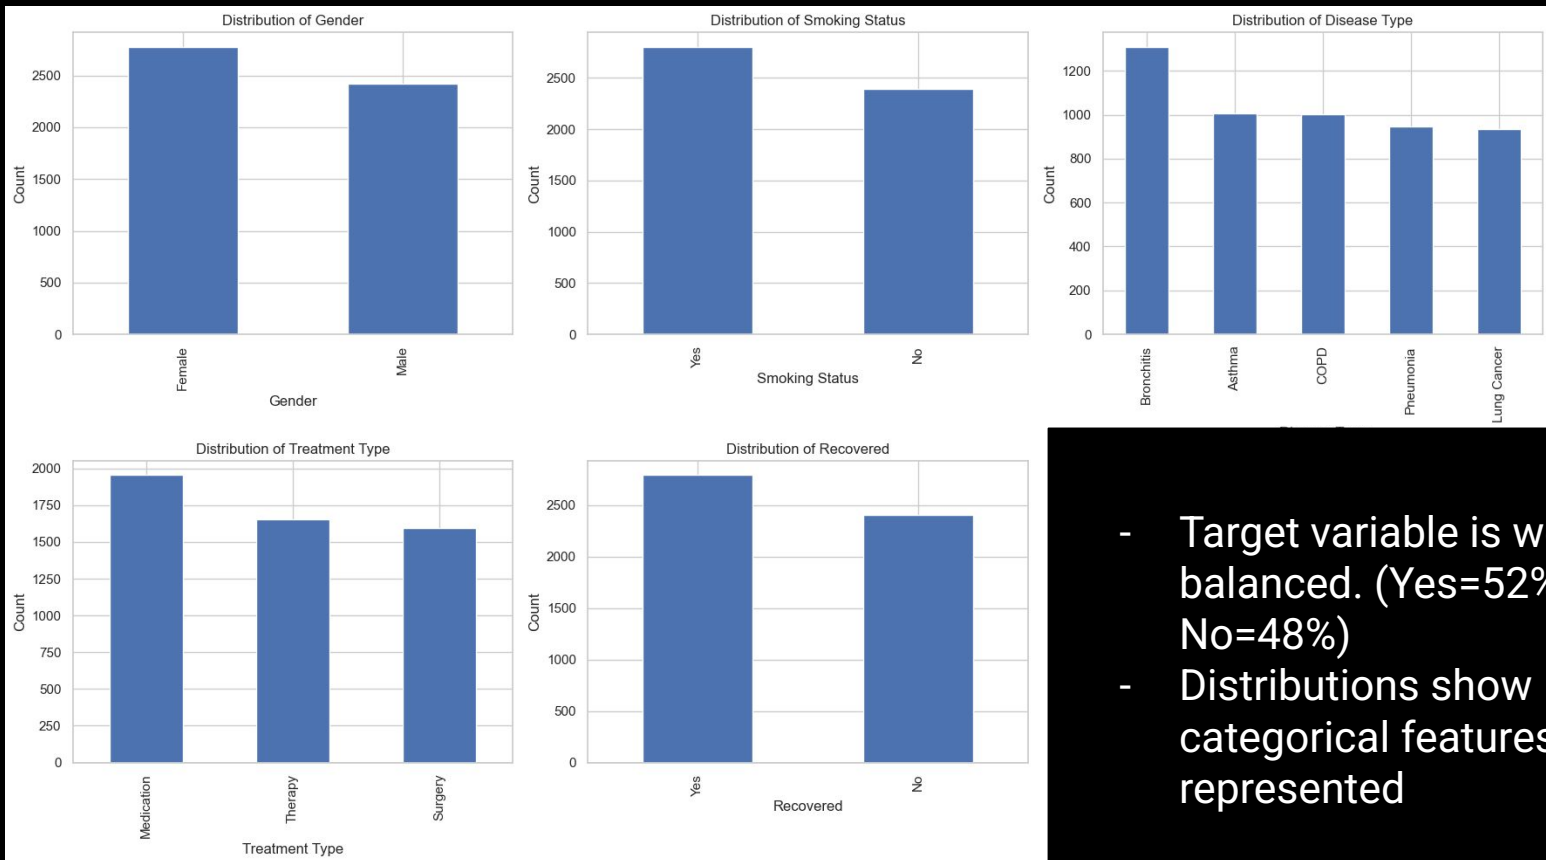
- Categorical: Gender, Smoking Status, Disease Type, Treatment Type
- Numerical: Age, Lung Capacity, Hospital Visits

Target Variable:

- Recovered: Categorical, if a patient recovered or not

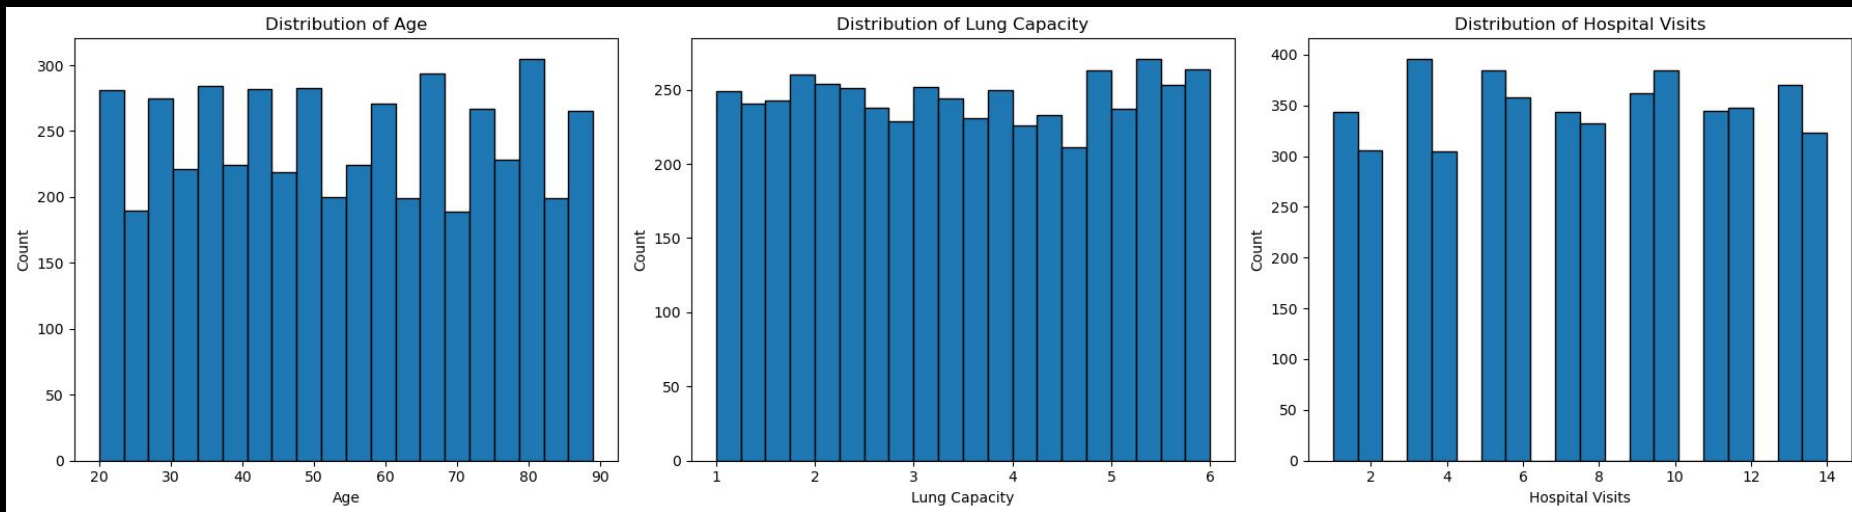
5200 Rows, 8 Columns

Exploratory Data Analysis



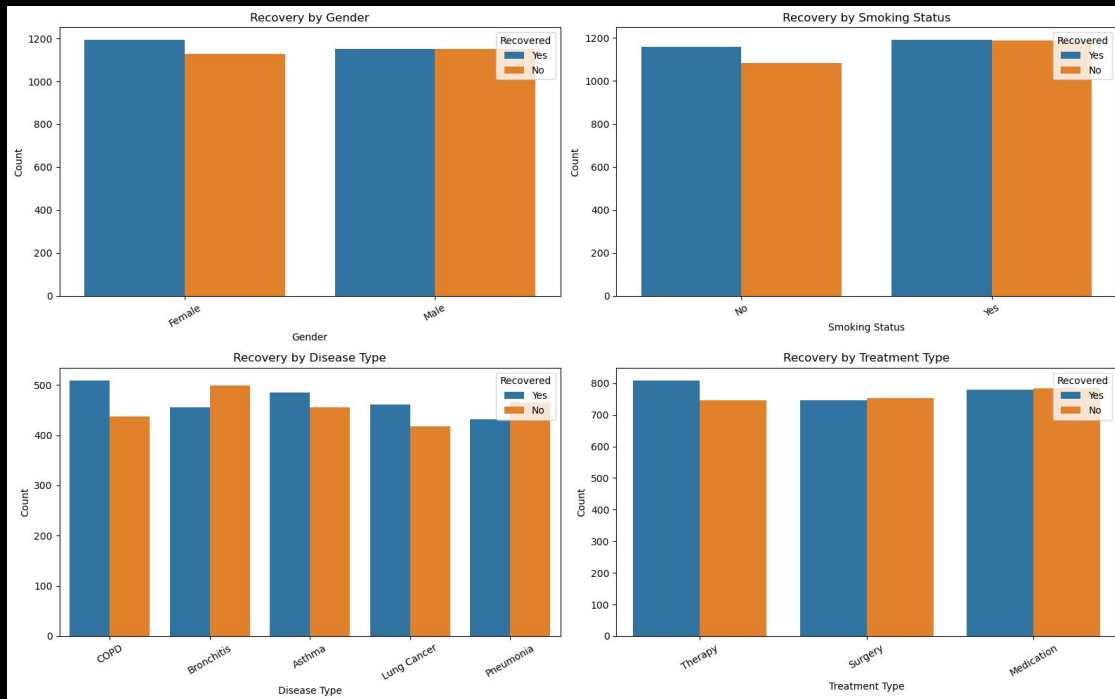
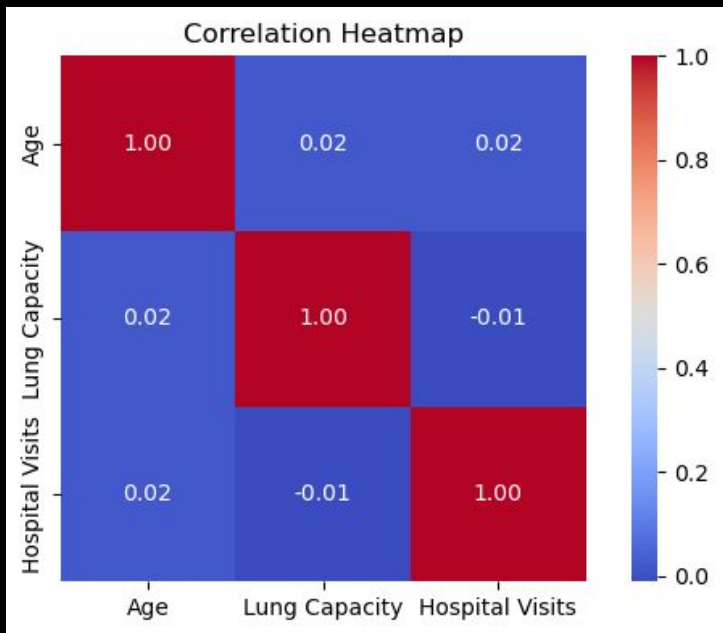
- Target variable is well balanced. (Yes=52%, No=48%)
- Distributions show categorical features are well represented

Exploratory Data Analysis



Numerical features also appear evenly distributed and well represented in the dataset.

Exploratory Data Analysis



Preprocessing:

```
#target variable encoding
df['Recovered']=df['Recovered'].map({'Yes': 1, 'No': 0})
#grouping features by there data types
cat_cols=df.select_dtypes(include='object').columns.tolist()
num_cols=df.select_dtypes(include=['int64', 'float64']).drop('Recovered', axis=1).columns.tolist()
#encoding categorical variables and scaling numerical variables
encoded = pd.DataFrame(OneHotEncoder(drop='first', sparse=False).fit_transform(df[cat_cols]))
scaled = pd.DataFrame(StandardScaler().fit_transform(df[num_cols]))
#recombining processed data
df_model = pd.concat([scaled, encoded], axis=1)
df_model['Recovered'] = df['Recovered'].values
```

Splitting:

Training and test splits were 20/80

Model Building

```
models={
    'Logistic Regression': LogisticRegression(max_iter=1000),
    'Random Forest': RandomForestClassifier(random_state=42),
    'XGBoost': XGBClassifier(use_label_encoder=False)}

for name, model in models.items():
    model.fit(X_train.values, y_train.values)
    y_pred = model.predict(X_test.values)
    y_train_pred = model.predict(X_train.values)
    cv_scores = cross_val_score(model, X_train.values, y_train.values, cv=5)
```



Trained each model on full training set and evaluated with cross validation.

Model Comparisons Pre-Tuning

Poor performances all around, barely above random guessing (50%)

Logistic Regression CV accuracy: mean = 0.5182, std = 0.0084

Logistic Regression Training classification report:

	precision	recall	f1-score	support
0	0.52	0.12	0.20	1904
1	0.54	0.90	0.67	2168
accuracy			0.54	4072
macro avg	0.53	0.51	0.44	4072
weighted avg	0.53	0.54	0.45	4072

Logistic Regression Test classification report:

	precision	recall	f1-score	support
0	0.55	0.13	0.21	452
1	0.57	0.92	0.70	567
accuracy			0.57	1019
macro avg	0.56	0.52	0.46	1019
weighted avg	0.56	0.57	0.48	1019

Random Forest CV accuracy: mean = 0.5174, std = 0.0068

Random Forest Training classification report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1904
1	1.00	1.00	1.00	2168
accuracy			1.00	4072
macro avg	1.00	1.00	1.00	4072
weighted avg	1.00	1.00	1.00	4072

Random Forest Test classification report:

	precision	recall	f1-score	support
0	0.46	0.41	0.44	452
1	0.57	0.62	0.59	567
accuracy			0.53	1019
macro avg	0.52	0.52	0.52	1019
weighted avg	0.52	0.53	0.52	1019

XGBoost CV accuracy: mean = 0.5066, std = 0.0058

XGBoost Training classification report:

	precision	recall	f1-score	support
0	0.94	0.90	0.92	1904
1	0.92	0.95	0.93	2168
accuracy			0.93	4072
macro avg	0.93	0.93	0.93	4072
weighted avg	0.93	0.93	0.93	4072

XGBoost Test classification report:

	precision	recall	f1-score	support
0	0.47	0.43	0.45	452
1	0.58	0.61	0.59	567
accuracy			0.53	1019
macro avg	0.52	0.52	0.52	1019
weighted avg	0.53	0.53	0.53	1019

Model Hyperparameter Tuning

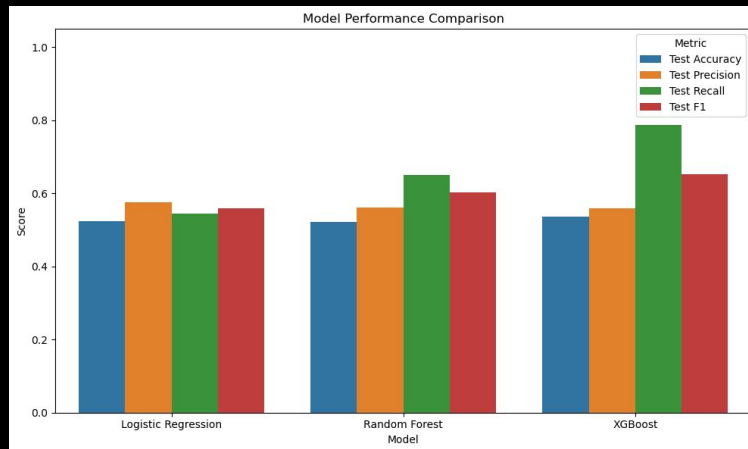
Tuned each model with
GridSearchCV:

- Only slight improvements, nothing significant
- LR: F1 & Precision improved
- RF: F1 & Recall improved
- XGBoost: F1 & Recall improved

Summary Table:

	Model	Train Accuracy	Test Accuracy	Train Precision
0	Logistic Regression	0.519892	0.524043	0.552180
1	Random Forest	0.995334	0.522080	0.991762
2	XGBoost	0.613212	0.535819	0.594277

	Test Precision	Train Recall	Test Recall	Train F1	Test F1
0	0.576493	0.519834	0.544974	0.535519	0.560290
1	0.560790	0.999539	0.650794	0.995635	0.602449
2	0.558897	0.862085	0.786596	0.703557	0.653480



Conclusions/Lessons Learned



- Each model had similar accuracies and performed poorly
- XGBoost performed the best overall (Highest F1-score and recall)
- Important to evaluate all scoring metrics to get full picture
- The dataset likely lacks strong predictive features and/or contains overlapping, noisy input variables.

End

THANK YOU